

Wistia Video Analytics (Data Engineering Academy)

1. Introduction and Business Goals

- **Scope:**
 - Designing the entire system architecture: ingestion, storage, processing, and reporting
 - Authenticating and ingesting data from the Wistia Stats API (both media and visitor level)
 - Handling pagination and incremental data pulls
 - Running the pipeline in production mode for 7 days
 - Implementing CI/CD using GitHub.
 - Documenting decisions, assumptions, and tradeoffs
- **Business Objectives:** The marketing team uses Wistia to track video engagement across Facebook and YouTube. We aim to:
 - Collect media-level and visitor-level analytics from Wistia's Stats API
 - Build an automated data pipeline to ingest and analyze performance
 - Use these insights to improve marketing strategies

This project simulates a real-world data engineering assignment with full responsibility placed on the student team to design, implement, and operate the system.

2. Requirements

This section details the specific requirements the solution must meet.

- You must design the architecture yourself and present it to SME for approval
- DO NOT USE DBT FOR TRANSFORMATION.
- DO NOT USE ANYTHING APART FROM AWS/Azure.
- You must use GitHub for version control and CI/CD
- You must use Python for API ingestion and PySpark for data transformation.

The pipeline should run for 7 days.

Functional Requirements:

- Design your own architecture for ingestion, processing, and storage

- Authenticate to Wistia Stats API using token-based Basic Auth
 - Extract media metadata (title, ID, hashed_id, created_at, etc.)
 - Extract engagement metrics (plays, play rate, watch time, etc.)
 - Extract visitor-level data (IP, engagement events)
 - Implement pagination to fetch all pages of results
 - Implement incremental ingestion based on created_at/updated_at
 - Run this pipeline in "production mode" for 7 consecutive days
 - Implement a CI/CD pipeline using GitHub Actions or equivalent
 - Store results in a structured data model (DWH or cloud database)
 - Create final reports or dashboards for insights (optional)
 - Submit a GitHub repo with documentation, pipeline code, CI/CD setup, and instructions
-

3. Architectural Design

This is the high-level overview of the system.

High-Level Diagram:



Technology Stack

Data Source

Wistia API

Architecture

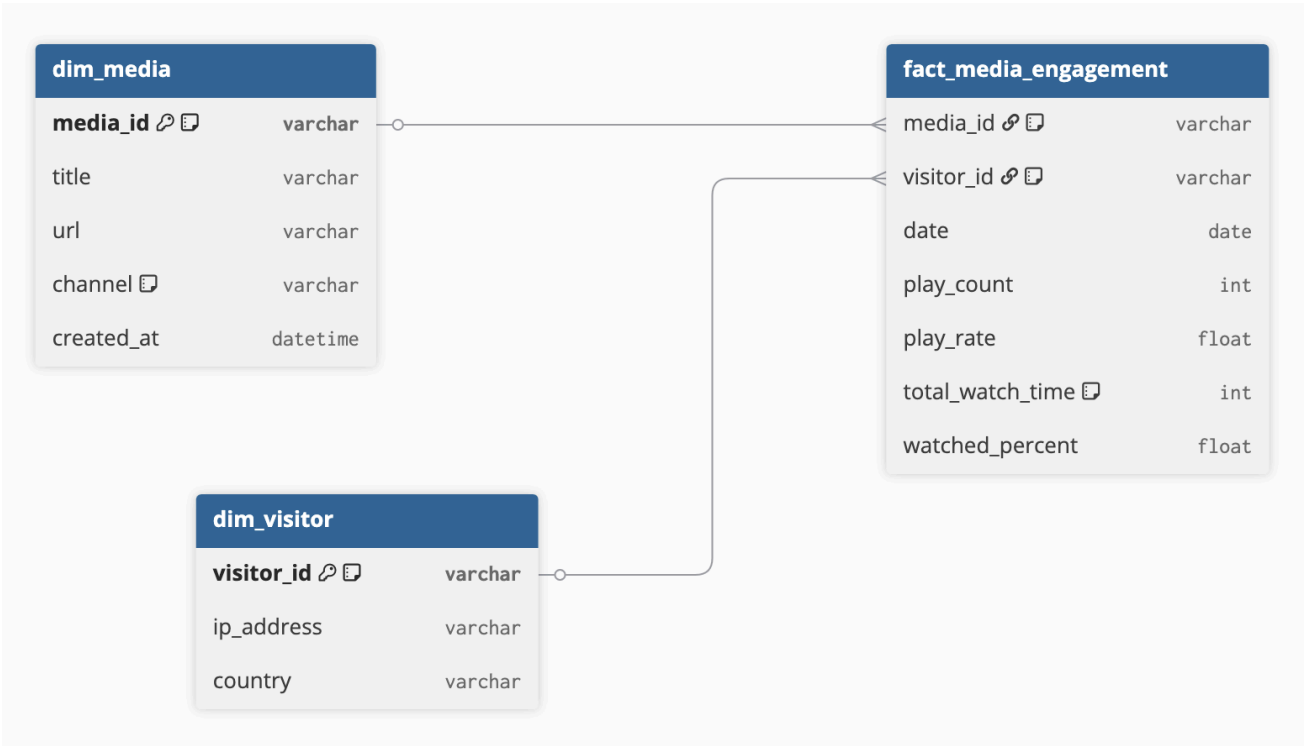
Bronze Stage

Lambda → S3

Dashboards (Streamlit) Then we can hit S3 directly with Streamlit on the views to create BI dashboards. (Optionally)

4. Data Model

- dim_media → https://docs.wistia.com/reference/get_stats-medias-mediaid
- dim_visitor → https://docs.wistia.com/reference/get_stats-visitors and
- fct_media_engagement → https://docs.wistia.com/reference/get_stats-medias-mediaid-engagement



Deliverables

Area	Criteria
Architecture	Clear, scalable, modular design
Data Quality	Correct use of pagination, incremental logic, and schema
Engineering	Effective error handling, retries, logging
CI/CD	Working CI/CD for deployment or validation
Documentation	README + architecture diagram + setup instructions